

Uncertainty Estimation with Noisy Optimizers

Master Seminar Beyond Deep Learning: Selected Topics on Novel Challenges

Orhun Güley

09.12.2020

Technical University of Munich

1. Introduction: Bayesian Inference
2. Variational Inference(VI)
3. Natural Gradient
4. Variational Inference using Noisy Natural Gradients¹
 - Noisy Adam
 - Noisy K-FAC
5. Variational Online Gauss-Newton(VOGN)²
 - Vadam
6. Comparison
7. Evaluation

¹Noisy Natural Gradient as Variational Inference (Zhang et al. [2018])

²Fast and Scalable Bayesian Deep Learning by Weight-Perturbation in Adam(Khan et al. [2018])

Introduction: Bayesian Inference

- Important tasks that have high risk or more interpretability require more than just "high accuracies".
- Autonomous driving, medical diagnosis, financial predictions for algorithmic trading...
- Bayesian Inference provides a framework where you can measure your confidence in your prediction.

Given a dataset $\mathcal{D} = \{(\mathbf{x}_i, y_i)_{i=1}^n\}$,

$$p(w|\mathcal{D}) = \frac{p(\mathcal{D}|w)p(w)}{p(\mathcal{D})} = \frac{p(\mathcal{D}|w)p(w)}{\int p(\mathcal{D}|w)p(w)} = \frac{p(\mathcal{D}|w)p(w)}{\sum_{i=1}^n p(\mathcal{D}|w)p(w)} \quad (1)$$

- **Frequentist** way: Directly maximizing the likelihood of the data given weights.
- In **Bayesian inference**, we put a prior distribution on weights and we try to find the posterior distribution
- **Likelihood**: $p(\mathcal{D}|w)$
- **Prior**: $p(w)$
- **Evidence**: $p(\mathcal{D}) = \int p(\mathcal{D}|w)p(w) \rightarrow$ Not tractable in complex models like deep neural networks!

- Due to the **intractability of the evidence term**, we do not have an exact solution. But we can approximate to posterior!
- **Several methods** → variational inference(VI), Markov Chain Monte Carlo(MCMC), and ensemble methods like MC-Dropout...
- In the recent years, it has been shown that with slight modifications on well-known optimizers, they perform variational inference!
- This presentation's topic is on variational inference.

Variational Inference(VI)

- Takes roots from relative entropy in information theory.
- A **non-negative, non-symmetric** measure between two probability distribution, defined as

$$\mathbb{KL}(q(x) \| p(x)) = \int q(x) \log \frac{q(x)}{p(x)} dx = \sum_{x \in X} p(x) \log \frac{p(x)}{q(x)}$$

- The integral version represents the continuous case and the summation version represents the discrete case

- We try to approximate the true posterior distribution with a variational distribution $q_{\theta}(\mathbf{w})$, which is parametrized by θ - the parameters of the posterior distribution.
- **Example:** If we model our approximate posterior q to be a Gaussian distribution

$$q_{\theta}(\mathbf{w}) := \mathcal{N}(w|\mu, \text{diag}(\sigma^2)) \quad \text{or} \quad q_{\theta}(\mathbf{w}) := \mathcal{N}(w|\mu, \Sigma)$$

, the parameters would be

$$\theta = (\mu, \sigma) \quad \text{or} \quad \theta = (\mu, \Sigma)$$

, where $\mu, \sigma \in \mathbb{R}^D$, and $\Sigma \in \mathbb{R}^{D \times D}$.

We try to maximize variational parameters θ :

$$\begin{aligned}\theta^* &= \arg \min_{\theta} \text{KL}[q_{\theta}(\mathbf{w}) \| p(\mathbf{w} | \mathcal{D})] \\ &= \arg \min_{\theta} \underbrace{\text{KL}[q_{\theta}(\mathbf{w} | \theta) \| p(\mathbf{w})] - \mathbb{E}_{q(\mathbf{w} | \theta)}[\log p(\mathcal{D} | \mathbf{w})]}_{\mathcal{F}(\mathcal{D}, \theta) = -\mathcal{L}(\theta, q) \rightarrow \text{negative ELBO}}\end{aligned}\tag{2}$$

It turns out that minimizing the KL-divergence is equal to maximizing the variational objective ELBO!

Reparametrization Trick:

In order to free our gradients of μ and Σ from the Gaussian distribution by using a parameter-free noise ϵ .

1. Draw samples from noise $\epsilon \sim \mathcal{N}(\mathbf{0}, I)$
2. Set $\mathbf{w} = t(\epsilon, \theta = \{\mu, \Sigma\}) = \mu + \Sigma \circ \epsilon$

Natural Gradient

- Natural gradient method takes its roots from information geometry.
- In optimization of the deep learning cost functions, learning takes place in the space of parameters, which is actually a space of probability distributions.
- Aims to change the gradients direction from the steepest direction in the Euclidean space to the steepest direction the **natural parameter space**.
- **Good news:** This space can be a probability distribution space!

Natural Gradient




Figure 1: Gradient and vs natural gradient directions for mean of variational distribution q . VI with a diagonal covariance is applied to the posterior $p(x, y) \propto \exp[-9(xy - 1)^2 - x^2 - y^2]$.¹




Figure 2: Arrows correspond the gradients and the curves correspond the path of gradients.²

¹A gradient-based algorithm competitive with variational bayesian em for mixture of gaussians, 2009 Kuusela et al. [2009]

²Natural gradients inpractice: Non-conjugate variational inference in gaussian process models Salimbeni et al. 9 / 29

- Natural gradient method scales the gradient with a inverse of **Fisher Information Matrix(FIM)** of a **score function**.
- In **probabilistic setting**, this score function is a **probability distribution**.
- The likelihood function can be a proper choice in this task, but in a Bayesian setting that we are trying to approximate the true posterior, our variational posterior $q_{\theta}(\mathbf{w})$ would be the right choice for score function
- FIM is given by:

$$F = \mathbb{E}_{q_{\theta}(\mathbf{w})} [\nabla_{\theta} q_{\theta}(\mathbf{w}) \nabla_{\theta} q_{\theta}(\mathbf{w})^T] \quad (3)$$


- Let $\tilde{\nabla}_\theta$ be the term for natural gradient. Then, the mathematical equation for the natural gradient of our variational objective is defined as

$$\begin{aligned}\tilde{\nabla}_\theta \mathcal{L}(\boldsymbol{\theta}, q) &\stackrel{\text{def}}{=} \nabla_\theta \mathcal{L}(\boldsymbol{\theta}) \mathbb{E}_z \left[(\nabla \log p_\theta(z))^T (\nabla \log p_\theta(z)) \right]^{-1} \\ &\stackrel{\text{def}}{=} \nabla_\theta \mathcal{L}(\boldsymbol{\theta}, q) \mathbf{F}^{-1}\end{aligned}\tag{4}$$

2 major variants of natural gradient that are frequently in use. The first one is named **natural gradient for point estimation(NGPE)** and the second one is **natural gradient for variational inference(NGVI)**:

1. **NGPE:** NGPE tries to estimate just the likelihood $p(\mathcal{D}|\mathbf{w})$ by optimizing over a loss function. For this task, the FIM is constructed by calculating the
$$\mathbf{F} = \text{Cov}_{\mathcal{D} \sim p(\mathcal{D})}[\nabla_{\mathbf{w}} \log p(\mathcal{D}|\mathbf{w})].$$
2. **NGVI:** We try to fit the parameters of a variational posterior $q_{\theta}(\mathbf{w})$ to maximize our variational objective ELBO. We need to compute the FIM of our variational posterior q instead of the predictive distribution, which is
$$\mathbf{F} = \text{Cov}_{\mathbf{w} \sim q_{\theta}(\mathbf{w})}[\nabla_{\mathbf{w}} \log q_{\theta}(\mathbf{w})].$$

Variational Inference using Noisy Natural Gradients^a



^aNoisy Natural Gradient as Variational Inference (Zhang et al. [2018])

The following updates for μ and Λ :

$$\begin{aligned}\mu &\leftarrow \mu + \alpha \Lambda^{-1} [\nabla_{\mathbf{w}} \log p(\mathcal{D}|\mathbf{w}) - \frac{\eta}{N} \mathbf{w}] \\ \Lambda &\leftarrow \left(1 - \frac{\beta}{N}\right) \Lambda - \beta [\nabla_{\mathbf{w}}^2 \log p(\mathcal{D}|\mathbf{w}), \mathbf{w}) - \frac{\eta}{N} \mathbf{I}]\end{aligned}\tag{5}$$

where $q_{\theta}(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mu, \Sigma)$, the precision matrix $\Lambda = \Sigma^{-1}$, a spherical Gaussian prior $p(\mathbf{w}) = \mathcal{N}(\mathbf{0}, \mathbf{I}/\eta)$, and then $\nabla_{\mathbf{w}}^2 \log p(\mathbf{w}) = \eta \mathbf{I}$.

Problems:

1. Hessian is hard to compute and not differentiable in every point(Ex: ReLU).
2. Hessian might have negative eigen-values if the negative log-likelihood is not convex, which means that our update for Λ might not be positive semi-definite.

Idea: Approximate Hessian with NGPE FIM $F = \text{Cov}_{\mathcal{D} \sim p(\mathcal{D})}[\nabla_{\mathbf{w}} \log p(\mathcal{D}|\mathbf{w})]$, which prevents both of the problems stated:

$$\Lambda \leftarrow \left(1 - \frac{\beta}{N}\right) \Lambda + \beta \left[\underbrace{\nabla_{\mathbf{w}} \log p(\mathcal{D}|\mathbf{w}) \nabla_{\mathbf{w}} \log p(\mathcal{D}|\mathbf{w})^T}_{F} + \frac{\eta}{N} \mathbf{I} \right] \quad (6)$$

Observation: With a fixed prior variance η , Λ will become a damped version of the moving average of the FIM

Rewriting the equations with the FIM approximation of Hessian:

$$\Lambda = N\bar{\mathbf{F}} + \eta\mathbf{I}$$

$$\bar{\mathbf{F}} \leftarrow (1 - \tilde{\beta})\bar{\mathbf{F}} + \tilde{\beta}[\nabla_{\mathbf{w}} \log p(\mathcal{D}|\mathbf{w}) \nabla_{\mathbf{w}} \log p(\mathcal{D}|\mathbf{w})^T] \quad (7)$$

$$\mu \leftarrow \mu + \tilde{\alpha} \left(\bar{\mathbf{F}} + \frac{\eta}{N} \mathbf{I} \right)^{-1} \left[\nabla_{\mathbf{w}} \log p(\mathcal{D}|\mathbf{w}) - \frac{\eta}{N} \mathbf{w} \right]$$

- Estimating Λ^{-1} as a full covariance Gaussian is unrealistic since the number of parameters needed for a full covariance matrix is $(\text{dim}(\mathbf{w}))^2$.
- Zhang et al. (2018) proposed a solution by approximating the FIM with a **diagonal matrix** $\bar{\mathbf{f}}$. For their noisy natural gradient version of Adam, the updates for μ and $\bar{\mathbf{f}}$ are:

$$\begin{aligned}\mu &\leftarrow \mu + \tilde{\alpha} \left[\nabla_{\mathbf{w}} \log p(\mathcal{D}|\mathbf{w}) - \frac{\eta}{N} \mathbf{w} \right] / \left(\bar{\mathbf{f}} + \frac{\eta}{N} \right) \\ \bar{\mathbf{f}} &\leftarrow (1 - \tilde{\beta}) \bar{\mathbf{f}} + \tilde{\beta} (\nabla_{\mathbf{w}} \log p(\mathcal{D}|\mathbf{w}))^2\end{aligned}\tag{8}$$

They use a Kronecker-factored approximation¹ to the FIM to perform efficient approximate natural gradient updates, given by

$$\begin{aligned}\mathbf{F}_l &= \mathbb{E} \left[\text{vec} \{ \mathcal{D}\mathbf{W}_l \} \text{vec} \{ \mathcal{D}\mathbf{W}_l \}^\top \right] = \mathbb{E} \left[\mathbf{g}_l \mathbf{g}_l^\top \otimes \mathbf{a}_l \mathbf{a}_l^\top \right] \\ &\approx \mathbb{E} \left[\mathbf{g}_l \mathbf{g}_l^\top \right] \otimes \mathbb{E} \left[\mathbf{a}_l \mathbf{a}_l^\top \right] = \mathbf{S}_l \otimes \mathbf{A}_l = \tilde{\mathbf{F}}_l\end{aligned}\tag{9}$$

where, \mathbf{a}_l and \mathbf{s}_l correspond to input activations and outputs at layer l , $\mathcal{D}\mathbf{v} = \nabla_{\mathbf{v}} \log p(\mathcal{D}|\mathbf{w})$, $\mathbf{g}_l = \mathcal{D}\mathbf{s}_l$ and the gradient of weights at layer l is $\mathcal{D}\mathbf{W}_l = \mathbf{a}_l \mathbf{g}_l^\top$.

Then we have the inverse FIM as

$$\begin{aligned}\tilde{\mathbf{F}}_l^{-1} \text{vec} \{ \nabla_{\mathbf{W}_l} h \} &= \mathbf{S}_l^{-1} \otimes \mathbf{A}_l^{-1} \text{vec} \{ \nabla_{\mathbf{W}_l} h \} \\ &= \text{vec} \left[\mathbf{A}_l^{-1} \nabla_{\mathbf{W}_l} h \mathbf{S}_l^{-1} \right]\end{aligned}\tag{10}$$

¹Optimizing Neural Networks with Kronecker-factored Approximate Curvature

- By plugging in this inverse FIM approximation to eqn. (7), we will have the MVG posterior. The update rule for $\bar{\mathbf{A}}_l$ and $\bar{\mathbf{S}}_l$ is as follows:

$$\begin{aligned}\bar{\mathbf{A}}_l &\leftarrow (1 - \tilde{\beta})\bar{\mathbf{A}}_l + \tilde{\beta}\mathbf{a}_l\mathbf{a}_l^\top \\ \bar{\mathbf{S}}_l &\leftarrow (1 - \tilde{\beta})\bar{\mathbf{S}}_l + \tilde{\beta}\mathcal{D}\mathbf{s}_l\mathcal{D}\mathbf{s}_l^\top\end{aligned}\tag{11}$$

- And the Σ_l can be decomposed as the Kronecker product of two terms:

$$\begin{aligned}\Sigma_l &= \frac{1}{N} [\mathbf{S}_l^\gamma]^{-1} \otimes [\mathbf{A}_l^\gamma]^{-1} \\ &\triangleq \frac{1}{N} \left(\bar{\mathbf{S}}_l + \frac{1}{\pi_l} \sqrt{\frac{\eta}{N}} \mathbf{I} \right)^{-1} \otimes \left(\bar{\mathbf{A}}_l + \pi_l \sqrt{\frac{\eta}{N}} \mathbf{I} \right)^{-1}\end{aligned}\tag{12}$$

Algorithm 1: Noisy K-FAC. Subscript l denotes layers, . We assume zero momentum for simplicity. Differences from standard K-FAC are shown in red.

```

input    :  $\alpha$ : Stepsize
input    :  $\beta$ : Exponential moving average parameter
input    :  $\lambda, \eta, \gamma_{\text{ex}}$ : KL weighting, prior variance, extrinsic damping term
input    : Stats and inverse update intervals  $T_{\text{stats}}$  and  $T_{\text{inv}}$ 
1  $k \leftarrow 0$  and initialize  $\{\boldsymbol{\mu}_l\}_{l=1}^L, \{\mathbf{S}_l\}_{l=1}^L, \{\mathbf{A}_l\}_{l=1}^L$  Calculate the intrinsic damping term  $\gamma_{\text{in}} = \frac{\eta}{N}$ , total damping term
    $\gamma = \gamma_{\text{in}} + \gamma_{\text{ex}}$ .
2 while not converged do
3    $k \leftarrow k + 1$ 
4    $\mathbf{W}_l \sim \mathcal{MN} \left( \mathbf{M}_l, \frac{1}{N} [\mathbf{A}_l^{\gamma_{\text{in}}}]^{-1}, [\mathbf{S}_l^{\gamma_{\text{in}}}]^{-1} \right)$ 
5   if  $k \equiv 0 \pmod{T_{\text{stats}}}$  then
6     Update the factors  $\{\mathbf{S}_l\}_{l=1}^L, \{\mathbf{A}_l\}_{l=0}^{L-1}$  using eq.(11)
7   if  $k \equiv 0 \pmod{T_{\text{inv}}}$  then
8     Calculate the inverses  $\{[\mathbf{S}_l^{\gamma}]^{-1}\}_{l=1}^L, \{[\mathbf{A}_l^{\gamma}]^{-1}\}_{l=0}^{L-1}$  using eq. (13).
9    $\mathbf{V}_l = \nabla_{\mathbf{W}_l} \log p(\mathcal{D}|\mathbf{W}) - \gamma_{\text{in}} \cdot \mathbf{W}_l$ 
10   $\mathbf{M}_l \leftarrow \mathbf{M}_l + \alpha [\mathbf{A}_l^{\gamma}]^{-1} \mathbf{V}_l [\mathbf{S}_l^{\gamma}]^{-1}$ 

```


Variational Online Gauss-Newton(VOGN)^a

^aFast and Scalable Bayesian Deep Learning by Weight-Perturbation in Adam(Khan et al. [2018])

Shown¹ that we can express the NGVI update in terms of the MLE objective.

- Let's denote the MLE objective $f_i(\boldsymbol{\theta}) := -\log p(\mathcal{D}_i | \boldsymbol{\theta})$ and minibatch stochastic-gradient estimates as

$$f(\boldsymbol{\theta}) := \frac{1}{N} \sum_{i=1}^N f_i(\boldsymbol{\theta}), \quad \hat{\mathbf{g}}(\boldsymbol{\theta}) := \frac{1}{M} \sum_{i \in \mathcal{M}} \nabla_{\boldsymbol{\theta}} f_i(\boldsymbol{\theta}), \quad (13)$$

- show the the NGVI update can be written in terms of stochastically approximated Hessian of f as

$$\boldsymbol{\mu}_{t+1} = \boldsymbol{\mu}_t - \beta_t (\hat{\mathbf{g}}(\boldsymbol{\theta}_t) + \tilde{\eta} \boldsymbol{\mu}_t) / (\mathbf{s}_{t+1} + \tilde{\eta}), \quad \mathbf{s}_{t+1} = (1 - \beta_t) \mathbf{s}_t + \beta_t \text{diag} \left[\hat{\nabla}_{\boldsymbol{\theta}\boldsymbol{\theta}}^2 f(\boldsymbol{\theta}_t) \right] \quad (14)$$

where $\boldsymbol{\theta}_t \sim \mathcal{N}(\boldsymbol{\theta} | \boldsymbol{\mu}_t, \boldsymbol{\sigma}_t^2)$ with $\boldsymbol{\sigma}_t^2 := 1/[N(\mathbf{s}_t + \tilde{\eta})]$ and $\tilde{\eta} := \eta/N$.

¹Fast and Scalable Bayesian Deep Learning by Weight-Perturbation in Adam(Khan et al. [2018])

Resembles the online-Newton method, since s_t contains an online estimate of the diagonal of the Hessian.

Problem: Hessian can be negative again!

Solution

Approximate Hessian with

$$\nabla_{\theta_j \theta_j}^2 f(\boldsymbol{\theta}) \approx \frac{1}{M} \sum_{i \in \mathcal{M}} [\nabla_{\theta_j} f_i(\boldsymbol{\theta})]^2 := \hat{h}_j(\boldsymbol{\theta}) \quad (15)$$

which is guaranteed to be positive if initial σ^2 is positive!

As a result, the new update rule for s_t becomes


$$\mathbf{s}_{t+1} = (1 - \beta_t) \mathbf{s}_t + \beta_t \hat{\mathbf{h}}_j(\boldsymbol{\theta}) \quad (16)$$

Variational Online Gauss-Newton (VOGN)

Algorithm 1: Variational Online Gauss Newton (VOGN)

```

1: Initialise  $\mu_0, \mathbf{s}_0, \mathbf{m}_0$ .
2:  $N \leftarrow \rho N, \tilde{\delta} \leftarrow \tau \delta / N$ .
3: repeat
4:   Sample a minibatch  $\mathcal{M}$  of size  $M$ .
5:   Split  $\mathcal{M}$  into each GPU (local minibatch  $\mathcal{M}_{local}$ ).
6:   for each GPU in parallel do
7:     for  $k = 1, 2, \dots, K$  do
8:       Sample  $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ .
9:        $\mathbf{w}^{(k)} \leftarrow \mu + \epsilon \sigma$  with  $\sigma \leftarrow (1/(N(\mathbf{s} + \tilde{\delta} + \gamma)))^{1/2}$ .
10:      Compute  $\mathbf{g}_i^{(k)} \leftarrow \nabla_{\mathbf{w}} \ell(\mathbf{y}_i, \mathbf{f}_{\mathbf{w}^{(k)}}(\mathbf{x}_i)), \forall i \in \mathcal{M}_{local}$ 
        using the method described in Appendix B.
11:       $\hat{\mathbf{g}}_k \leftarrow \frac{1}{M} \sum_{i \in \mathcal{M}_{local}} \mathbf{g}_i^{(k)}$ .
12:       $\hat{\mathbf{h}}_k \leftarrow \frac{1}{M} \sum_{i \in \mathcal{M}_{local}} (\mathbf{g}_i^{(k)})^2$ .
13:    end for
14:     $\hat{\mathbf{g}} \leftarrow \frac{1}{K} \sum_{k=1}^K \hat{\mathbf{g}}_k$  and  $\hat{\mathbf{h}} \leftarrow \frac{1}{K} \sum_{k=1}^K \hat{\mathbf{h}}_k$ .
15:  end for
16:  AllReduce  $\hat{\mathbf{g}}, \hat{\mathbf{h}}$ .
17:   $\mathbf{m} \leftarrow \beta_1 \mathbf{m} + (\hat{\mathbf{g}} + \tilde{\delta} \mu)$ .
18:   $\mathbf{s} \leftarrow (1 - \tau \beta_2) \mathbf{s} + \beta_2 \hat{\mathbf{h}}$ .
19:   $\mu \leftarrow \mu - \alpha \mathbf{m} / (\mathbf{s} + \tilde{\delta} + \gamma)$ .
20: until stopping criterion is met
  
```



Learning rate	α
Momentum rate	β_1
Exp. moving average rate	β_2
Prior precision	δ
External damping factor	γ
Tempering parameter	τ
# MC samples for training	K
Data augmentation factor	ρ

Figure 3: Table is prepared from the experimental results of [Khan et al., 2018, Zhang et al., 2018]

- Assuming that variational posterior q is from an exponential-family distribution with natural parameter η , Khan et al. (2018) proposed the following natural-momentum method:

$$\eta_{t+1} = \eta_t + \bar{\alpha}_t \tilde{\nabla}_{\eta} \mathcal{L}_t + \bar{\gamma}_t (\eta_t - \eta_{t-1}) \quad (17)$$

where $\tilde{\nabla}$ is the natural-gradients with respect to natural parameter space and the gradient scaled by the Fisher information matrix of $q(\theta)$.

- They show that the eqn.(17) can be expressed as VON update with momentum:

$$\begin{aligned}\mu_{t+1} &= \mu_t - \bar{\alpha}_t \left[\frac{1}{\mathbf{s}_{t+1} + \tilde{\eta}} \right] \circ (\nabla_{\theta} f(\boldsymbol{\theta}_t) + \tilde{\eta} \boldsymbol{\mu}_t) + \bar{\gamma}_t \left[\frac{\mathbf{s}_t + \tilde{\eta}}{\mathbf{s}_{t+1} + \tilde{\eta}} \right] \circ (\boldsymbol{\mu}_t - \boldsymbol{\mu}_{t-1}) \\ \mathbf{s}_{t+1} &= (1 - \bar{\alpha}_t) \mathbf{s}_t + \bar{\alpha}_t \nabla_{\theta\theta}^2 f(\boldsymbol{\theta}_t)\end{aligned}\tag{18}$$

where $\mathbf{w}_t \sim \mathcal{N}(\boldsymbol{\theta} \mid \boldsymbol{\mu}_t, \sigma_t^2)$ with $\sigma_t^2 := 1/[N(\mathbf{s}_t + \tilde{\lambda})]$.

Noisy Adam vs Vadam

Algorithm 2: Noisy Adam. Differences from standard Adam are shown in red.

input : α : Stepsize
input : β_1, β_2 : Exponential decay rates for updating μ and \mathbf{f}
input : η, γ_{ex} : prior variance, extrinsic damping term

```
1  $\mathbf{m} \leftarrow 0$  ;  
2 Calculate the intrinsic damping term  $\gamma_{\text{in}} = \frac{\eta}{N}$ ,  
   total damping term  $\gamma = \gamma_{\text{in}} + \gamma_{\text{ex}}$ .  
3 while not converged do  
4    $\mathbf{w} \sim \mathcal{N}(\mu, \frac{1}{N} \text{diag}(\mathbf{f} + \gamma_{\text{in}})^{-1})$   
5    $\mathbf{g} \leftarrow \nabla_{\mathbf{w}} \log p(\mathcal{D}|\mathbf{w})$   
6    $\mathbf{m} \leftarrow \beta_1 \cdot \mathbf{m} + (1 - \beta_1) \cdot (\mathbf{g} + \gamma_{\text{in}} \cdot \mathbf{w})$   
7    $\mathbf{f} \leftarrow \beta_2 \cdot \mathbf{f} + (1 - \beta_2) \cdot (\mathbf{g} \circ \mathbf{g})$ .  
   (Update momentum)  
8    $\tilde{\mathbf{m}} \leftarrow \mathbf{m} / (1 - \beta_1^k)$   
9    $\hat{\mathbf{m}} \leftarrow \tilde{\mathbf{m}} / (\mathbf{f} + \gamma)$   
10   $\mu \leftarrow \mu + \alpha \cdot \hat{\mathbf{m}}$  (Update parameters)
```

Algorithm 3: Vadam. Differences from standard Adam are shown in red

input : α : Stepsize
input : β_1, β_2 : Exponential decay rates for updating μ and \mathbf{f}
input : η : prior variance

```
1  $\mathbf{m} \leftarrow 0$  ;  
2 while not converged do  
3    $\mathbf{w} \leftarrow \mu + \sigma \circ \epsilon$ , where  
    $\epsilon \sim \mathcal{N}(0, \mathbf{I}), \sigma \leftarrow 1/\sqrt{Ns + \lambda}$   
4    $\mathbf{g} \leftarrow -\nabla \log p(\mathcal{D}_i | \theta)$   
5    $\mathbf{m} \leftarrow \beta_1 \mathbf{m} + (1 - \beta_1) (\mathbf{g} + \mu \frac{\eta}{N})$   
6    $\mathbf{s} \leftarrow \beta_2 \mathbf{s} + (1 - \beta_2) (\mathbf{g} \circ \mathbf{g})$   
7    $\hat{\mathbf{m}} \leftarrow \mathbf{m} / (1 - \beta_1^t), \hat{\mathbf{s}} \leftarrow \mathbf{s} / (1 - \beta_2^t)$   
8    $\mu \leftarrow \mu - \alpha \hat{\mathbf{m}} / (\sqrt{\hat{\mathbf{s}}} + \frac{\eta}{N})$ 
```

Vadam and VOGN Convergence




Figure 4: Results on the Australian-Scale dataset using a neural network with a hidden layer of 64 units for different minibatch sizes M and number of MC samples S . Figure taken from Khan et al. [2018]

Comparison

Comparison of Optimizers

	Noisy Adam	Vadam	Noisy K-FAC	VOGN
Prior	Spherical Gaussian	Spherical Gaussian	Spherical Gaussian	Spherical Gaussian
Posterior	Gaussian with diag. cov.	Gaussian with diag. cov.	Matrix-Variate Gaussian	Gaussian with diag. cov.
FIM Approximation	GM	GM	GM	GGN
Gradient Estimators of Gaussian	Oppor Estimator	Reparametrization Trick	Oppor Estimator	Reparametrization Trick

Table 1: Difference and similarities between the noisy optimizers. GM corresponds to Gradient Magnitude, and GGN corresponds to Generalized Gauss-Newton approximation.

- **Noisy Adam and Vadam** \leftarrow algorithmically very similar and both easy to implement.
- Derivation of **Vadam** uses **natural-momentum** term from **Polyak's heavy ball method**, which noisy Adam doesn't provide.
- **Noisy K-FAC** \rightarrow more complex weight distribution due to covariances.

Evaluation

	Test log-likelihood				
Dataset	MC-Dropout	BBVI	Vadam	Noisy Adam	Noisy K-FAC
Boston	-2.46 +/- 0.06	-2.73 +/- 0.05	-2.85 +/- 0.07	-2.558 +/- 0.032	-2.417 +/- 0.029
Concrete	-3.04 +/- 0.02	-3.24 +/- 0.02	-3.39 +/- 0.02	-3.145 +/- 0.023	-3.039 +/- 0.025
Energy	-1.99 +/- 0.02	-2.47 +/- 0.02	-2.15 +/- 0.07	-1.629 +/- 0.020	-1.421 +/- 0.005
Kin8nmm	0.95 +/- 0.01	0.95 +/- 0.01	0.76 +/- 0.00	1.112 +/- 0.008	1.148 +/- 0.007
Naval	3.80 +/- 0.01	4.46 +/- 0.03	4.72 +/- 0.22	6.231 +/- 0.041	7.079 +/- 0.034
Power	-2.80 +/- 0.01	-2.88 +/- 0.01	-2.88 +/- 0.01	-2.803 +/- 0.010	-2.776 +/- 0.011
Wine	-0.93 +/- 0.01	1.00 +/- 0.001	-1.00 +/- 0.01	-0.976 +/- 0.016	-0.969 +/- 0.014
Yacht	-1.55 +/- 0.03	-2.41 +/- 0.02	-1.70 +/- 0.03	-2.412 +/- 0.006	-2.316 +/- 0.006

Table 2: Comparison of Noisy Adam, Noisy K-FAC and Vadam with other popular methods. Table is prepared from the experimental results of [Khan et al., 2018, Zhang et al., 2018]

Dataset/ Architecture	Optimiser	Train/Validation Accuracy (%)	Validation NLL	Epochs	Time/ epoch (s)	ECE	AUROC
CIFAR-10/ LeNet-5 (no DA)	Adam	71.98 / 67.67	0.937	210	6.96	0.021	0.794
	BBB	66.84 / 64.61	1.018	800	11.43 [†]	0.045	0.784
	MC-dropout	68.41 / 67.65	0.99	210	6.95	0.087	0.797
	VOGN	70.79 / 67.32	0.938	210	18.33	0.046	0.8
CIFAR-10/ AlexNet (no DA)	Adam	100.0 / 67.94	2.83	161	3.12	0.262	0.793
	MC-dropout	97.56 / 72.20	1.077	160	3.25	0.140	0.818
	VOGN	79.07 / 69.03	0.93	160	9.98	0.024	0.796
CIFAR-10/ AlexNet	Adam	97.92 / 73.59	1.480	161	3.08	0.262	0.793
	MC-dropout	80.65 / 77.04	0.667	160	3.20	0.114	0.828
	VOGN	81.15 / 75.48	0.703	160	10.02	0.016	0.832
CIFAR-10/ ResNet-18	Adam	97.74 / 86.00	0.55	160	11.97	0.082	0.877
	MC-dropout	88.23 / 82.85	0.51	161	12.51	0.166	0.768
	VOGN	91.62 / 84.27	0.477	161	53.14	0.040	0.876
ImageNet/ ResNet-18	SGD	82.63 / 67.79	1.38	90	44.13	0.067	0.856
	Adam	80.96 / 66.39	1.44	90	44.40	0.064	0.855
	MC-dropout	72.96 / 65.64	1.43	90	45.86	0.012	0.856
	OGN	85.33 / 65.76	1.60	90	63.13	0.128	0.854
	VOGN	73.87 / 67.38	1.37	90	76.04	0.029	0.854
	K-FAC	83.73 / 66.58	1.493	60	133.69	0.158	0.842
	Noisy K-FAC	72.28 / 66.44	1.44	60	179.27	0.080	0.852

Figure 5: Performance comparisons on different dataset/architecture combinations. Figure taken from Osawa et al. [2019]

- The distributed version of momentum employed VOGN method proposed by Osawa et al. (2019) has very promising results as well.
- VOGN and Noisy K-FAC has a comparable performance with standart Adam and MC-Dropout if data augmentation(DA) is applied.

Questions?

References

- Mohammad Emtiyaz Khan, Didrik Nielsen, Voot Tangkaratt, Wu Lin, Yarin Gal, and Akash Srivastava. Fast and scalable bayesian deep learning by weight-perturbation in adam, 2018.
- Mikael Kuusela, T. Raiko, Antti Honkela, and J. Karhunen. A gradient-based algorithm competitive with variational bayesian em for mixture of gaussians. *2009 International Joint Conference on Neural Networks*, pages 1688–1695, 2009.
- Kazuki Osawa, Siddharth Swaroop, Anirudh Jain, Runa Eschenhagen, Richard E. Turner, Rio Yokota, and Mohammad Emtiyaz Khan. Practical deep learning with bayesian principles, 2019.
- Hugh Salimbeni, Stefanos Eleftheriadis, and James Hensman. Natural gradients in practice: Non-conjugate variational inference in gaussian process models, 2018.

Guodong Zhang, Shengyang Sun, David Duvenaud, and Roger Grosse. Noisy natural gradient as variational inference, 2018.